Inferring the existence and direction of causal associations in the face of measurement error

01 February 2017

Gibran Hemani*, Kate Tilling and George Davey Smith

MRC Integrative Epidemiology Unit (IEU) at the University of Bristol, School of Social and Community Medicine, Bristol, UK

* Correspondence to: g.hemani@bristol.ac.uk

Abstract

Inference of the causal relationship between two traits can be achieved by combining genetic associations with a mediation-based approach, as is done in the causal inference test (CIT) and others. We show here, however, that measurement error in the phenotypes can lead to mediation-based approaches inferring the wrong causal direction, and that increasing sample sizes has the adverse effect of increasing confidence in the wrong answer. Here we develop an extension to Mendelian randomisation, a method that uses genetic associations in an instrumentation framework, that enables inference of the causal direction with some advantages. First, it is less susceptible to bias in the presence of measurement error; second, it is more statistically efficient; and third, it can be performed using only summary level data from genome-wide association studies. We apply the method to infer the causal direction between DNA methylation and gene expression levels. Our results demonstrate that, in general, DNA methylation is more likely to be the causal factor, but that such high level conclusions are susceptible to differential measurement error between the platforms. We emphasise that the use of several different methods is important to triangulate reliable conclusions.

Introduction

Observational measures of the human phenome are growing ever more abundant, but using these data to make causal inference is notoriously susceptible to many pitfalls, with basic regression-based techniques unable to distinguish a true causal association from reverse causation or confounding[1–3]. In response to this, genetic instrumentation has emerged as a technique for improving the reliability of causal inference in observational data, and with the coincident rise in genome-wide association studies it is now a prominent tool that is applied in several different guises[3–6]. However, potential pitfalls remain and one that is often neglected is the influence of non-differential measurement error on the reliability of causal inference.

Measurement error (or observational error, henceforth referred to simply as measurement error) is the difference between the measured value of a quantity and its true value. This study focuses specifically on non-differential measurement error where all strata of a measured variable have the same error rate, which can manifest as changes in scale or levels of imprecision (noise) of the measurement. Such variability can arise through a whole plethora of mechanisms, which are often unique to the study design and difficult to avoid[7,8]. Array technology is now commonly used to obtain high throughput phenotyping at low cost, but comes with the problem of having imperfect resolution, for instance methylation levels as measured by the Illumina450k chip are prone to have some amount of noise around the true value due to imperfect sensitivity[9,10]. Relatedly, if the measurement of biological interest is the methylation level in a T cell, then measurement error of this value can be introduced by using methylation levels from whole blood samples because the measured value will be an assay of many cell types[11].

Measurement error will of course arise in more basic data too, for example when measuring BMI one is typically interested in using this as a proxy for adiposity, but it is clear that the correlation between BMI and underlying adiposity is not perfect[12]. A similar problem of biological misspecification is unavoidable in

disease diagnosis, and measuring behaviour such as smoking or diet is notoriously difficult to do accurately. Measurement error can also be introduced after the data has been collected, for example the transformation of non-normal data for the purpose of statistical analysis will lead to a new variable that will typically incur both changes in scale and imprecision (noise) compared to the original variable. The sources of measurement error are not limited to this list[8], and its impact has been explored in the epidemiological literature extensively[13,14]. Given the near-ubiquitous presence of measurement error in phenomic data it is vital to understand its impact on the tools we use for causal inference.

An established study design that can provide information about causality is randomisation. Given the hypothesis that trait A (henceforth referred to as the exposure) is causally related to trait B (henceforth referred to as the outcome), randomisation can be employed to assess the causal nature of the association by randomly splitting the sample into two groups, subjecting one group to the exposure and leaving the other as a control. The association between the exposure and the outcome in this setting provides a robust estimate of the causal relationship. This provides the theoretical basis behind randomised control trials, but in practice randomisation is often impossible to implement in an experimental context due to cost, scale or inability to manipulate the exposure. The principle, however, can be employed in extant observational data through the use of genetic variants associated with the exposure (instruments), where the inheritance of an allele serves as a random lifetime allocation of differential exposure levels[15,16]. Two statistical approaches to exploiting the properties of genetic instruments are widely used: mediation-based approaches and Mendelian randomisation (MR).

Mediation-based approaches employ genetic instruments (typically single nucleotide polymorphisms, SNPs) to orient the causal direction between the exposure and the outcome. If a SNP is associated with an exposure, and the exposure is associated with some outcome, then it logically follows that the indirect influence of the SNP on the outcome will disappear when conditioning on the exposure. Here, the exposure 'mediates' the association between the SNP and the outcome, providing information about the causal influence of the exposure on the outcome. This forms the basis of a number of methods such as genetical genomics[17], the regression-based causal inference test (CIT)[4,18], a structural equation modelling (SEM) implementation in the NEO software[5], and various other methods including Bayesian approaches[6]. They have been employed by a number of recent publications that make causal inferences in large scale 'omics datasets[6,19–23].

MR can be applied to the same data - phenotypic measures of the exposure and the outcome variables and a genetic instrument for the exposure - but the genetic instrument is employed in a subtly different manner. Here the SNP is used as a surrogate for the exposure. Assuming the SNP associates with the outcome only through the exposure, the causal effect of the exposure on the outcome can be estimated by scaling the association between the SNP and the outcome by the association between the SNP and the exposure. Though difficult to test empirically, this assumption can be relaxed in various ways when multiple instruments are available for a putative exposure[24,25] and a number of sensitivity tests are now available to improve reliability[26].

By utilising genetic instruments in different ways, mediation-based analysis and MR models have properties that confer some advantages and some disadvantages for reliable causal inference. An important advantage of mediation-based analysis is that the true underlying biology of the genetic instrument doesn't necessarily need to be known as long as it associates with the putative exposure. In the CIT framework (described fully in the Methods) for example, the test statistic is different if you test for the exposure causing the outcome or the outcome causing the exposure, allowing the researcher to infer the direction of causality between two variables by performing the test in both directions and choosing the model with the strongest evidence. The CIT can also distinguish between several putative models of association between the traits and the SNP[4,18]. Such is not the case for MR, where in order to infer the direction of causality between two traits the instrument must be known to influence the exposure primarily, associating with the outcome only through the exposure.

Assuming biological knowledge of genetic associations can be problematic because if there exists a putative association between two variables, with the instrumental SNP being robustly associated with each, it can be difficult to determine which of the two variables is subject to the primary effect of the SNP (i.e. for which of the two variables is the SNP a valid instrument? Figure 1). By definition, we expect that if the association is causal then a SNP for the exposure will be associated with the outcome, such that if the researcher

erroneously uses the SNP as an instrument for the outcome then they are likely to see an apparently robust causal association of outcome on exposure. Such a situation can arise in many scenarios. Genome-wide association studies (GWASs) that identify genetic associations for complex traits are, by design, hypothesis free and agnostic of genomic function, and it often takes years of follow up studies to understand the biological nature of a putative GWAS hit[27]. If multiple instruments are available for an hypothesised exposure, which is increasingly typical for complex traits that are analysed in large GWAS consortia, then techniques can be applied to mitigate these issues[16]. But these techniques cannot always be applied in the case of determining causal directions between 'omic measures where typically only one cis-acting SNP is known. For example if a DNA methylation probe is associated with expression of an adjacent gene, then is a cis-acting SNP an instrument for the DNA methylation level, or the gene expression level[28] (Figure 1)?

MR does however have some important advantages over the mediation-based approaches. The mediation-based approaches require that the exposure, outcome and instrumental variables are all measured in the same data, whereas recent extensions to MR circumvent this requirement, allowing causal inference to be drawn when exposure variables and outcome variables are measured in different samples[29]. This has the crucial advantage of improving statistical power by allowing analysis in much larger sample sizes, and dramatically expands the breadth of possible phenotypic relationships that can be evaluated[26]. Additionally, when MR assumptions are satisfied the method is substantially more robust to there being measurement error in the exposure variable. Indeed instrumental variable (IV) analysis was in part initially introduced as a correction for measurement error in the exposure[30], whereas it has been noted that both classic mediation-based analyses[13,14,31,32] and mediation-based methods that use instrumental variables[33,34] are prone to be unreliable in its presence.

Using theory and simulations we show how non-differential measurement error in phenotypes can lead to unreliable causal inference in the mediation-based CIT method. We then present an extension to MR that allows researchers to ascertain the causal direction of an association even when the biology of the instruments are not fully understood, and also a method to evaluate the sensitivity of the result of this extension to measurement error. This improves the utility of MR in cases where mediation based methods might have otherwise been used preferentially. Finally, we apply this new method to infer the direction of causation between DNA methylation levels and gene expression levels.

Methods

The analysis proceeds by simulating variables that have known causal effects and valid instruments, and then applying measurement error in different configurations. We then assess the performance of the CIT test and MR analysis for a) inference of a causal relationship and b) inferring the correct direction of causality. Included in the MR analysis is a simple extension that is used to make inference of the causal direction when it is unknown which of the correlated variables the instrument is directly acting upon. All analysis was performed using the R programming language[35] and code is made available at github url to go here and implemented in the MR-Base (http://wwww.mrbase.org) platform[26].

CIT test

The CIT method[4] is implemented in the R package R/cit[18]. The methodology of the CIT is as follows. Assume an exposure x is instrumented by a SNP g, and the exposure x is causally related to an outcome y. Thus,

$$g \sim Binom(2, 0.5)$$
$$x = \alpha_g + \beta_g g + \epsilon_g$$
$$y = \alpha_x + \beta_x x + \epsilon_x$$

The following tests are then performed:

1. $H_0 : cov(g, x) = 0; H_1 : cov(g, x) \neq 0;$ *the SNP associates with the exposure*

3

2. $H_0 : cov(g, y) = 0; H_1 : cov(g, y) \neq 0$; *the SNP associates with the outcome*
3. $H_0 : cov(x, y) = 0; H_1 : cov(x, y) \neq 0$; *the exposure associates with the outcome*
4. $H_0 : cov(g, y - \hat{y}) \neq 0; cov(g, y - \hat{y}) = 0$; *the SNP is independent of the outcome when the outcome is adjusted for the exposure*

where $y - \hat{y} = y - \hat{\alpha}_g + \hat{\beta}_g x$ is the residual of y after adjusting for x, where x is assumed to mediate the association between the SNP and the outcome. If all four tests reject the null hypothesis then it is inferred that x is causally related to y. The CIT measures the strength of causality by generating an omnibus p-value, p_{CIT}, which is simply the largest (least extreme) p-value of the four tests, the intuition being that causal inference is only as strong as the weakest link in the chain of tests.

In these analyses the *cit.cp* function was used to obtain an omnibus p-value. To infer the direction of causality using the CIT method, an omnibus p-value generated by CIT, p_{CIT}, was estimated for the direction of x causing y, and for the direction model of y causing x. For some significance threshold α, the existence of causality and its direction was inferred based on the following scenarios:

- If $p_{CIT,x \to y} < \alpha$ and $p_{CIT,y \to x} > \alpha$ then the $x \to y$ model is accepted
- If $p_{CIT,x \to y} > \alpha$ and $p_{CIT,y \to x} < \alpha$ then the $y \to x$ model is accepted
- If $p_{CIT,x \to y} > \alpha$ and $p_{CIT,y \to x} > \alpha$ then no inference is made
- If $p_{CIT,x \to y} < \alpha$ and $p_{CIT,y \to x} < \alpha$ then the confounding model is accepted ($x \leftarrow g \to y$).

For the purposes of compiling simulation results we use an arbitrary $\alpha = 0.05$ value, though we stress that for real analyses it is not good practice to rely on p-values for making causal inference, nor is it reliable to depend on arbitrary significance thresholds[36].

MR causal test

Two stage least squares (2SLS) is a commonly used technique for performing MR when the exposure, outcome and instrument data are all available in the same sample. Assuming the following model

$$y = \alpha_{MR} + \beta_{MR}\hat{x} + \epsilon_{MR}$$

where $\hat{x} = \hat{\alpha}_g + \hat{\beta}_g g$, a causal relationship is inferred if the null hypothesis that $\beta_{MR} = 0$ is rejected. A p-value for this test, p_{MR}, was obtained using the R package $R/systemfit$[37].

The method that we will now describe is able to distinguish between two models, $x \to y$ or $y \to x$. Unlike the CIT framework, this approach cannot infer if the true model is $x \leftarrow g \to y$. We also assume all genetic effects are additive.

To infer the direction of causality it is desirable to know which of the variables, x or y, is being directly influenced by the instrument g. This can be achieved by assessing which of the two variables has the biggest absolute correlation with g (Appendix 2), formalised by testing for a difference in the correlations ρ_{gx} and ρ_{gy} using Steiger's Z-test for correlated correlations within a population[38]. It is calculated as

$$Z = (Z_{gx} - Z_{gy}) \frac{\sqrt{N - 3}}{\sqrt{2(1 - \rho_{xy})h}}$$

where Fisher's z-transformation is used to obtain $Z_{g*} = \frac{1}{2} \ln\left(\frac{1+\rho_{g*}}{1-\rho_{g*}}\right)$,

$$h = \frac{1 - (frm^2)}{1 - rm^2}$$

where

$$f = \frac{1 - \rho_{xy}}{2(1 - rm^2)}$$

and

$$rm^2 = \frac{1}{2}(\rho_{gx}^2 + \rho_{gy}^2).$$

The Z value is interpreted such that

$$Z \begin{cases} > 0, & x \to y \\ < 0, & y \to x \\ = 0, & x \perp\!\!\!\perp y \end{cases}$$

and a p-value is generated from the Z value to indicate the probability of obtaining the observed difference in correlations ρ_{gx} and ρ_{gy} under the null hypothesis that both correlations are identical. The existence of causality and its direction is inferred based on the following scenarios:

- If $p_{Steiger} < \alpha$ and $p_{MR} < \alpha$ and $Z > 0$ then a causal association for the correct model is accepted
- If $p_{Steiger} < \alpha$ and $p_{MR} < \alpha$ and $Z < 0$ then a causal association for the incorrect model is accepted
- Otherwise no inference is made

For the purposes of compiling simulation results, we use an arbitrary $\alpha = 0.05$ value.

Note that the same correlation test approach can be applied to a two-sample MR[29] setting. Two-sample MR refers to the case where the SNP-exposure association and SNP-outcome association are calculated in different samples (e.g. from publicly available summary statistics). Here the Steiger test of two independent correlations can be applied where.

$$Z = \frac{Z_{gx} - Z_{gy}}{\sqrt{1/(N_1 - 3) + 1/(N_2 - 3)}}$$

An advantage of using the Steiger test in the two sample context is that it can compare correlations in independent samples where sample sizes are different. Steiger test statistics were calculated using the *r.test* function in the R package *R/psych*[39].

Steiger sensitivity analysis

The Steiger test for inferring if $x \to y$ is based on evaluating $\rho_{gx} > \rho_{gy}$. However, ρ_{gx} and ρ_{gy} are underestimated if x and y are measured imprecisely. If, for example, x has lower measurement precision than y then we might empirically obtain $\rho_{g,x_o} > \rho_{g,y_o}$. As we show in the appendix, the maximum measurement error of x_o is ρ_{g,x_o} and the minimum is 0 (same applies to y_o). It is possible to simulate what the inferred causal direction would be for all values within these bounds. To evaluate how reliable, R, the Steiger test is to potential measurement imprecision in x and y we integrate over the simulated R values for the possible measurement error values which are compatible with the direction of causality obtained empirically from the Steiger test. We also do the same for the measurement error values which are not compatible. The ratio of these quantities is an indication of how much of the measurement error space supports the inferred causal direction. A ratio $R = 1$ indicates that the Steiger test is highly sensitive to measurement error, with equal weight of the measurement error parameter space supporting both directions of causality. Full details are provided in Appendix 2.

Simulations

Simulations were conducted by creating variables of sample size n for the exposure x, the measured values of the exposure x_o, the outcome y, the measured values of the outcome y_o and the instrument g. One of two models are simulated, the "causal model" where x causes y and g is an instrument for x; or the "non-causal model" where g influences a confounder u which in turn causes both x and y. Here x and y are correlated but not causally related. Each variable in the causal model was simulated such that:

$$g \sim Binom(2, 0.5)$$
$$x = \alpha_g + \beta_g g + \epsilon_g$$
$$x_o = \alpha_{mx} + \beta_{mx} x + \epsilon_{mx}$$
$$y = \alpha_x + \beta_x x + \epsilon_x$$
$$y_o = \alpha_{my} + \beta_{my} y + \epsilon_{my}$$

where $\epsilon_{m*} \sim N(0, \sigma^2_{m*})$, α_{mx} and β_{mx} are parameters that represent non-differential measurement error into the exposure variable x, and α_{my} and β_{my} are parameters for non-differential measurement error in the outcome y. Similarly in the non-causal model:

$$g \sim Binom(2, 0.5)$$
$$u = \alpha_g + \beta_g g + \epsilon_g$$
$$x = \alpha_u + \beta_u u + \epsilon_{u_x}$$
$$x_o = \alpha_{mx} + \beta_{mx} x + \epsilon_{mx}$$
$$y = \alpha_u + \beta_u u + \epsilon_{u_y}$$
$$y_o = \alpha_{my} + \beta_{my} y + \epsilon_{my}$$

All α values were set to 0, and β values set to 1. Normally distribted values of ϵ_* were generated such that

$$cor(g, x)^2 = 0.1$$
$$cor(g, u)^2 = 0.1$$
$$cor(x, y)^2 = \{0.2, 0.4, 0.6, 0.8\}$$
$$\sigma^2_{mx} = \{0, 0.2, 0.4, 0.6, 0.8, 1\}$$
$$\sigma^2_{my} = \{0, 0.2, 0.4, 0.6, 0.8, 1\}$$
$$n = \{100, 1000, 10000\}$$

giving a total of 432 combinations of parameters. Simulations using each of these sets of variables were performed 100 times, and the CIT and MR methods were applied to each in order to evaluate the causal association of the simulated variables. Similar patterns of results were obtained for different values of $cor(g, x)$ and $cor(g, u)$.

Two sample MR

Two sample MR[29] was performed using the summary statistics presented in[28] for each of 4122 associations between DNA methylation and gene expression levels. To do this we obtained a list of 458 gene expression - DNA methylation associations as reported in Shakhbazov et al[28]. These were filtered to be located on the same chromosome, have robust correlations after correcting for multiple testing, and to share a SNP that had a robust cis-acting effect on both the DNA methylation probe and the gene expression probe. Because only summary statistics were available (effect, standard error, effect allele, sample size, p-values) for the instrumental SNP on the methylation and gene expression levels, the Steiger test of two independent

correlations was used to infer the direction of causality for each of the associations. The Wald ratio test was then used to estimate the causal effect size for the estimated direction for each association.

Results

The influence of measurement error on mediation

Measurement error of an exposure can be modeled as some transformation of the true value that leads to the observed value, $x_o = f(x)$. For example, we can define $f(x) = \alpha_{mx} + \beta_{mx}x + \epsilon_{mx}$, where α_{mx} and β_{mx} influence the error in the measurement of x by altering its scale, and ϵ_{mx} represents the imprecision (or noise) in the measurement of x. Here the true value of the exposure is partially explained by the genetic instrument, g, such that

$$x = \alpha_g + \beta_g g + \epsilon_g$$

where β_g is the effect of the SNP on the exposure, and ϵ_g is the normally distributed residual value of x; and the outcome is partially explained by the exposure

$$y = \alpha_x + \beta_x x + \epsilon_x$$

where β_x is the true effect of the exposure on the outcome. In the causal inference test (CIT), an omnibus p-value is generated from four hypothesis tests: 1) g is associated with x; 2) g is associated with y; 3) x is associated with y; and 4) g is independent of $y|x$. The 4th condition employs mediation for causal inference, and can be expressed as $cov(g, y - \hat{y}) = 0$, where $\hat{y} = \hat{\alpha}_x + \hat{\beta}_x x_o$. When measurement error in scale and imprecision is introduced, such that y_o is the measured value of y, it can be shown using basic covariance properties (Appendix 1) that

$$cov(g, y - \hat{y}) = cov(g, y_o) - cov(g, \hat{y}_o)$$
$$= \beta_{my}\beta_g\beta_x var(g) - D\beta_{my}\beta_g\beta_x var(g)$$

where

$$D = \frac{\beta_{mx}^2 var(x)}{\beta_{mx}^2 var(x) + var(\epsilon_{mx})}$$

Thus an observational study will find $cov(g, y_o - \hat{y}_o) = 0$ when the true model is causal only when $D = 1$. Therefore, if there is any measurement error that incurs imprecision (i.e. $var(\epsilon_{mx}) \neq 0$) then there will remain an association between g and $y_o|x$, which is in violation of the the 4th condition of the CIT. Note that scale transformation of x or y without any incurred imprecision is insufficient to lead to a violation of the test statistic assumptions, and henceforth mention of measurement error will relate to imprecision unless otherwise stated.

We performed simulations to verify that this problem does arise using the CIT method. Figure 2 shows that when there is no measurement error in the exposure or outcome variables ($\rho_{x,x_o} = 1$) the CIT is reliable in identifying the correct causal direction. However, as measurement error increases in the exposure variable, eventually the CIT is more likely to infer a robust causal association in the wrong direction. Also of concern here is that increasing sample size does not solve the issue, indeed it only strengthens the apparent evidence for the incorrect inference.

Using MR to infer the existence of causality

When selecting an instrument g that has a direct influence on x where x is causally related to y, one can reason that the influence of g on y is the proportion of the effect of x on y that is explained by the effect of g on x. Hence the effect estimate β_{MR} of x on y is estimated as

$$\beta_{MR} = \frac{\beta_{y_o}}{\beta_g} = \frac{cov(y_o, g)}{cov(x_o, g)}$$
$$= \frac{\beta_{my}\beta_x cov(x, g)}{\beta_{mx} cov(x, g)}$$
$$= \frac{\beta_{my}}{\beta_{mx}}\beta_x$$

Thus we obtain an estimate of the effect of x_o on y that is scaled by the measurement error in x and y, but is unrelated to measurement imprecision. Measurement error and imprecision are likely to reduce the power of MR inferring a causal association (increased false negative rate), but importantly the evidence for causal association, based on rejecting the null hypothesis that $\beta_{MR} = 0$, is not inflated when the null hypothesis is true.

We performed simulations to compare the performance of MR against CIT in detecting a causal association between simulated variables under different levels of imprecision simulated in the exposure. Figure 4 shows the true positive rates between the CIT and MR for detecting a causal association. We observe that the CIT has lower power in all cases, with performance declining as measurement imprecision increases in the exposure. When MR assumptions are satisfied, notably that it is known on which of x and y the SNP g has a primary influence, the performance of MR in detecting an association is unrelated to measurement error in the exposure. Measurement error in the outcome does reduce power, but does not induce a substantive difference in performance between CIT and MR.

Using MR to infer the direction of causality

If we do not know whether the SNP g has a direct influence on x or y then further efforts are required. If x causes y and g is a valid instrument for x then we expect that $\rho_{g,x}$ to be greater than $\rho_{g,y}$ because the association of g with y is the proportion of the true causal association between x and y scaled by the association between g and x. We can use the Steiger test of correlated correlations to formally test the difference in magnitude of $\rho_{g,x}$ and $\rho_{g,y}$, where the variable with the largest correlation on g is deemed to be the variable which is imposing the causal effect.

This is a simple method to orient the direction of causality in an MR analysis when the underlying biology of the SNP is not fully understood. But it can be shown that in the presence of measurement imprecision, $d = \rho_{x,x_o} - \rho_{x,y}\rho_{y,y_o}$ (Appendix 2), where d is a metric of the degree to which the Steiger test is liable to provide the wrong estimate (*i.e.* if $d > 0$ then the Steiger test is likely to be correct about the causal direction). Thus it is important to note that under certain conditions when measurement imprecision is present the Steiger test could make erroneous inference about the causal direction because the value of d is not restricted to being greater than 0. However Figure 7 shows that when there is no measurement error in x, the Steiger test is unlikely to infer the wrong direction of causality even if there is measurement error in y. It also shows that in most cases where x is measured with error, especially when the causal effect between x and y is not very large, the sensitivity of the Steiger test to measurement error is relatively low.

We performed simulations to explore the performance of this approach in comparison to CIT in terms of the rate at which a strong causal relationship is obtained for the correct direction of causality, and the rate at which a strong causal relationship is obtained where the reported direction of causality is incorrect. Simulations were performed for two models, one for a "causal model" where there was a causal relationship between x and y; and one for a "non-causal model" where x and y were not causally related, but had a confounded association induced by the SNP g influencing a confounder variable u.

Figure 5a shows that, for the "causal model", when $d < 0$ the MR analysis is indeed liable to infer the wrong direction of causality, and that this erroneous result is more likely to occur with increasing sample size. However, the CIT is in general more fallable to reporting a robust causal association for the wrong direction of causality. When $d > 0$ is satisfied we observe that in most cases the MR method has greater power to obtain evidence for causality than CIT, and always obtains the correct direction of causality. The CIT, unlike the Steiger test for MR, is able to distinguish the "non-causal model" from the "causal model" (Methods, Figure 5b), but it is evident that measurement error will often lead the CIT to evaluate this confounding model as being causal.

MR analysis of gene expression and DNA methylation levels

We used the Steiger test to infer the direction of causality between DNA methylation levels and gene expression levels and found that the causal direction commonly goes in both directions (Figure 6a). We then went on to predict the causal directions of the associations for varying levels of possible measurement error. Figure 6a shows that the qualitative result of which of DNA methylation or gene expression is more commonly the causal factor is very strongly determined by measurement error.

We performed two sample MR[29] for each association in the direction of causality that we estimated to be the correct direction from the Steiger test. We observed that the sign of the MR estimate was generally in the same direction as the Pearson correlation coefficient reported by Shakhbazov et al[28], but with a few exceptions (Figure 6b). There was a moderate correlation between the absolute magnitudes of the causal correlation and the observational Pearson correlation (r = 0.4538942). Together these inferences suggest that even for relatively low level associations it is important to use causal inference methods over observational associations to infer causal effect sizes.

We also observed that for associations where methylation caused gene expression the causal effect was more likely to be negative than for the associations where gene expression caused methylation (OR = 0.6103961 (95% CI 0.3622428 - 1.029752), Figure 6c), suggesting that reducing gene expression levels at a controlling CpG typically leads to increased gene expression levels, consistent with expectation[40].

Discussion

Researchers are often confronted with the problem of making causal inferences using a statistical framework on observational data. In the epidemiological literature issues of measurement error in mediation analysis is relatively well explored[41]. Our analysis extends this to related methods such as CIT that are widely used in predominantly 'omic data. These methods are indeed liable to the same issues as standard mediation based analysis, and specifically we show that as measurement error in the exposure variable increases, CIT is likely to have reduced statistical power, and liable to infer the wrong direction of causality. We also demonstrate that, though unintuitive, increasing sample size does not resolve the issue, rather it leads to more extreme p-values for the model that predicts the wrong direction of causality.

Under many circumstances a practical solution to this problem is to use Mendelian randomisation instead of methods such as the CIT or similar that are based on mediation. Inferring the existence of causality using Mendelian randomisation is robust in the face of measurement error and, if the researcher has knowledge about the biology of the instrument being used in the analysis, can offer a direct solution to the issues that the CIT faces. This assumption is often reasonable, for example SNPs are commonly used as instruments when they are found in genes with known biological relevance for the trait of interest. But on many occasions this is not the case, and it may be tempting to resort to using methods based on mediation in order to be able to both ascertain if there is a causal association and to infer the direction of causality. Here we have described a simple extension to MR which can be used as an alternative to mediation based methods. We show that this method is still liable to measurement error, but because it has different properties to the CIT it offers two main advantages. First, it uses a formal statistical framework to test for the robustness of the assumed direction of causality. Second, after testing in a comprehensive range of scenarios the MR

9

based approach is less likely to infer the wrong direction of causality compared to CIT, while substantially improving power over CIT in the cases where $d > 0$.

We demonstrate this new method by evaluating the causal relationships of 458 known associations between DNA methylation and gene expression levels using summary level data. The causal direction is heavily influenced by how much measurement error is present in the different assaying platforms. For example, if DNA methylation measures typically have higher measuremet error than gene expression measures then our analysis suggests that DNA methylation levels would be more often the causal factor in the association. Indeed, previous studies which have evaluated measurement error in these platforms do support this position[42,43], though making strong conclusions for this analysis is difficult because measurement error is likely to be study specific.

In our simulations we focused on the simple case of a single instrument in a single sample setting with a view to making a fair comparison between MR and the various mediation-based methods available. However, MR strategies exist for other scenarios, for example if there are multiple instruments for the hypothesised exposure then the problems described here may be circumvented[16]. Likewise, if there is at least one instrument for each trait then bi-directional MR can offer solutions to inferring the causal direction[44]. In this work we assumed that pleiotropy (the influence of the instrument on the outcome through a mechanism other than the exposure) was not present. Recent method developments in MR[24,25] have focused on accounting for the issues that horizontal pleiotropy can introduce, but how they perform in the presence of measurement error remains to be explored. An important advantage that MR confers over most mediation based analysis is that it can be performed in two samples, which can considerably improve power and expand the scope of analysis. However, whether there is a substantive difference in two sample MR versus one sample MR in how measurement error has an effect is not yet fully understood. We have also assumed no measurement error in the genetic instrument, which is not unreasonable given the strict QC protocols that ensure high quality genotype data is available to most studies. We have restricted the scope to only exploring non-differential measurement error and avoided the complications incurred if measurement error in the exposure and outcome is correlated, however our analysis goes beyond many previous explorations of measurement error by assessing the impacts of both imprecision (noise) and linear transformations of the true variable on causal inference.

Mediation based network approaches, that go beyond analyses of two variables, are very well established[34] and have a number of extensions that make them valuable tools, including for example network construction. But because they are predicated on the basic underlying principles of mediation they are liable to suffer from the same issues of measurement error. Recent advances in MR methodology, for example applying MR to genetical genomics[45], multivariate MR[46] and mediation through MR[47,48] may offer more robust alternatives for these more complicated problems.

The overarching result from our simulations is that, regardless of the method used, inferring causal direction using an instrument of unknown biology is highly sensitive to measurement error. With the presence of measurement error near ubiquitous in most observational data, and our ability to measure it limited, we argue that it needs to be central to any consideration of approaches which are used in attempt to strengthen causal inference, and any putative results should be accompanied with appropriate sensitivity analysis that assesses their robustness under varying levels of measurement error.

Figures

Figure 1: Gene expression levels (blue blocks) and DNA methylation levels (green triangles) may be correlated but the causal direction is unknown. If a SNP (yellow circle) is associated with both DNA methylation and gene expression levels then it can be used as an instrument, but there are three competing models for these variables. The causal inference test (CIT) attempts to delineate between them. a) Methylation causes gene expression. The left figure shows that the SNP influences methylation levels that in turn influence gene expression levels. The right figure shows the directed acyclic graph that represents this model. Faded symbols represent the measured values whereas solid symbols represent the true values. b) The same as in A, except the causal direction is from gene expression to DNA methylation. c) A model of confounding, where gene expression and DNA methylation are not causally related, but the SNP influences them each through separate pathways or a confounder.

Figure 2: The CIT was performed on simulated variables where the exposure influenced the outcome and the exposure was instrumented by a SNP. The test statistic from CIT when testing if the exposure caused the outcome (the true model) is in red, and the test for the outcome causing the exposure (false model) is in green. Rows of plots represent the sample sizes used for the simulations. As measurement error increases (decreasing values on x-axis) the test statistic for the incorrect model gets stronger and the test statistic for the correct model gets weaker.

Figure 3: a) We can predict the values the Steiger test would take (z-axis) for different potential values of measurement error (x and y axes), drawn here as the blue surface. When $\rho_{g,y} > \rho_{g,x}$, as denoted by range of values where the blue surface is above the red plane, those values of measurement error lead to our observed Steiger test inferring the wrong causal direction. Where the blue surface lies below the red plane, these measurement error values support the inferred causal direction of X to Y. A measure of reliability, therefore, is the ratio of the negative and positive volumes of the total space bound by the blue and red surfaces, $R = \frac{V_{z \geq 0}}{-V_{z<0}}$. In this case, where $\rho_{g,x}^2 = 0.01$ and $\rho_{x,y}^2 = 0.1$, the $R = 4.40$, which means that 4.40 times as much of the possible measurement error values are in support of the $x \to y$ direction of causality than $y \to x$. b) The influence of $\rho_{x,y}$ (x-axis) on the Steiger reliability values (y-axis) for different values of $\rho_{g,x}^2$. Inferring the direction of causality is more sensitive to measurement error when there are stronger causal associations between X and Y.

Figure 4: Outcomes were simulated to be causally influenced by exposures with varying degrees of measurement imprecision applied to the exposure variable (x axis). True positive rates (y axis) for MR and CIT were compared for varying levels of measurement imprecision in the outcome variable (rows of boxes) and sample sizes (columns of boxes).
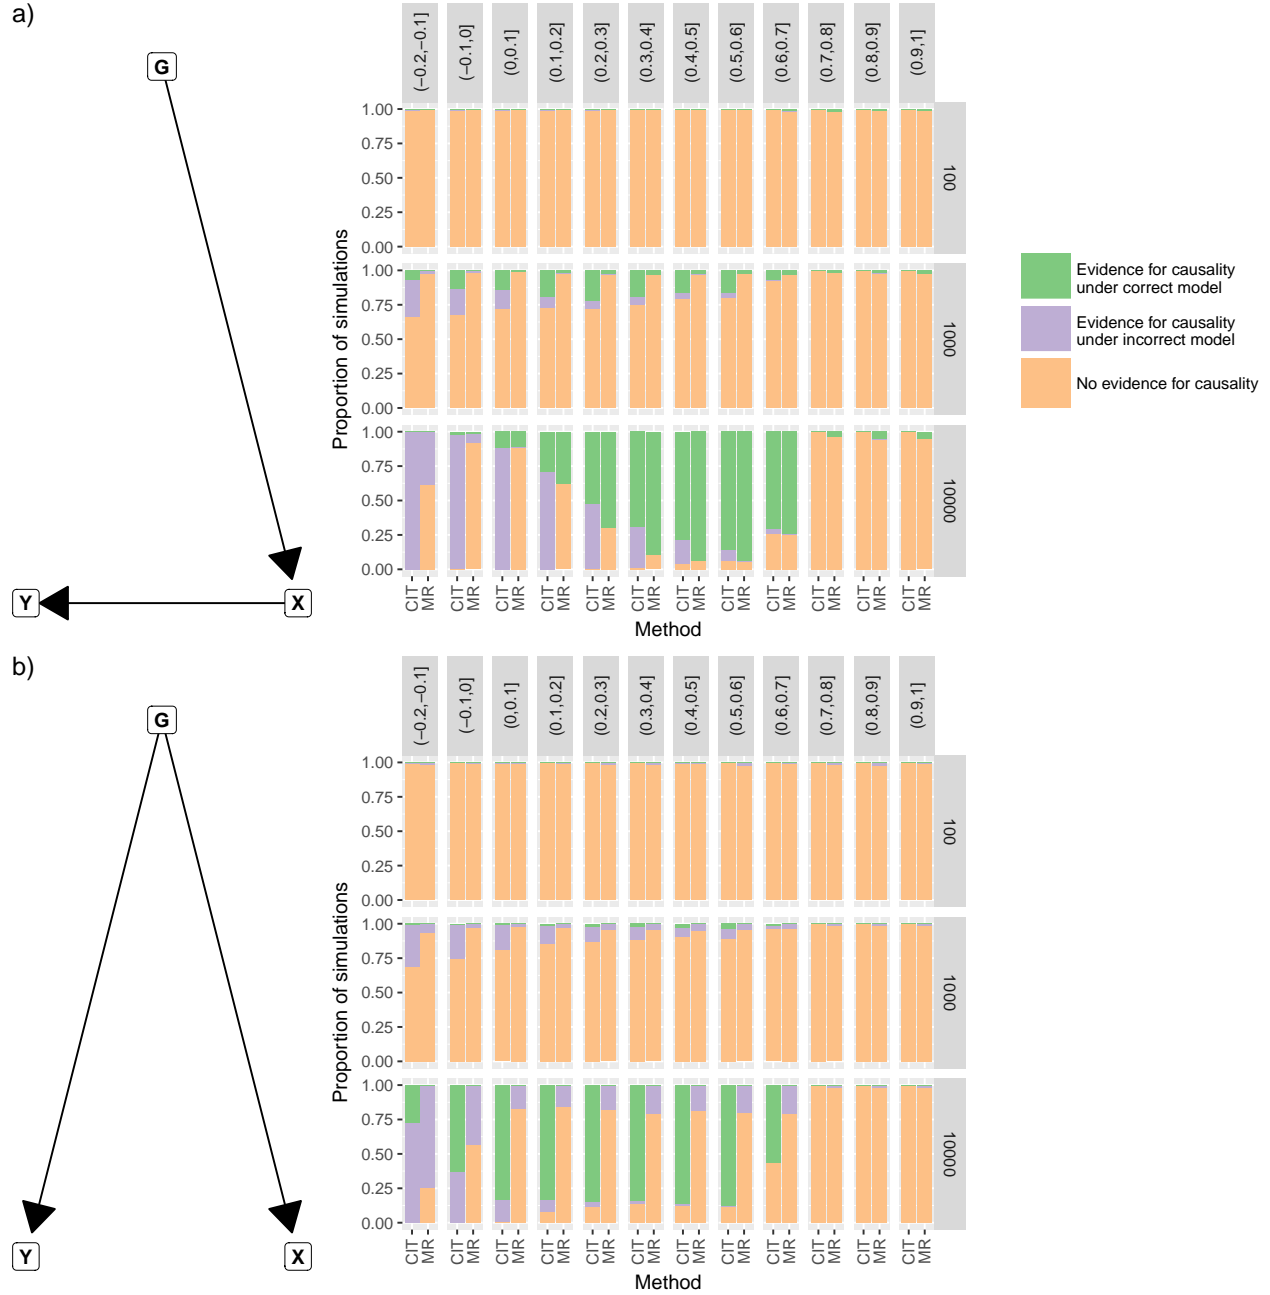
Figure 5: a) Outcomes were simulated to be causally influenced by exposures as shown in the graph, with varying degrees of measurement error applied to both. CIT and MR were used to infer evidence for causality between the exposure and outcome, and to infer the direction of causality. The value of $d = \rho_{x,x_o} - \rho_{x,y}\rho_{y,y_o}$, such that when d is negative we expect the Steiger test to be more likely to be wrong about the direction of causality. Rows of graphs represent the sample size used in the simulations. For the CIT method, outcome 1 denoted evidence for causality with correct model, outcomes 2 or 3 denoted evidence for causality with incorrect model, and outcome 4 denoted no evidence for causality. b) As in (a) except the simulated model was non-causal, and a genetic confounder induces an association between x and y. MR is unable to identify this model, so any significant associations are deemed to be incorrect. Outcome 3 denotes evidence for the correct model for the CIT method.

Figure 6: Using 458 putative associations between DNA methylation and gene expression we used the Steiger test to infer the direction of causality between them. a) The rightmost bar shows the proportion of associations for each of the two possible causal directions (colour key) assuming no measurement error in either gene expression or DNA methylation levels. The proportions change when we assume different levels of measurement error in gene expression levels (x-axis) or DNA methylation levels (columns of boxes). If there is systematically higher measurement error in one platform it will appear to be less likely to be the causal factor. b) The relationship between the Pearson correlation between DNA methylation and gene expression levels (x-axis) and the causal estimate (scaled to be in standard deviation units, y-axis). c) Distribution of estimated causal effect sizes, stratified into associations inferred to be due to DNA methylation causing expression (blue) and expression causing DNA methylation (red).

Appendix 1

We assume the following model

$$x = \alpha_g + \beta_g g + \epsilon_g$$
$$x_o = \alpha_{mx} + \beta_{mx} x + \epsilon_{mx}$$
$$y = \alpha_x + \beta_x x + \epsilon_x$$
$$y_o = \alpha_{my} + \beta_{my} y + \epsilon_{my}$$

where x is the exposure on the outcome y, g is an instrument that has a direct effect on x, x_o is the measured quantity of x, where measurement error is incurred from linear transformation in α_{mx} and β_{mx} and imprecision from ϵ_{mx}, and y_o is the measured quantity of y, where measurement error is incurred from linear transformation in α_{my} and β_{my} and imprecision from ϵ_{my}. Our objective is to estimate the expected magnitude of association between g and y after conditioning on x. Under the CIT, this is expected to be $cov(g, y_o - \hat{y}_o) = 0$ when x causes y, where $\hat{y}_o = \hat{a}_{x_o} + \hat{\beta}_{x_o} x_o$ is the predicted value of y_o using the measured value of x_o.

We can split $cov(g, y_o - \hat{y}_o)$ into two parts, $cov(g, y_o)$ and $cov(g, \hat{y}_o)$.

Part 1

$$
\begin{aligned}
cov(g, y_o) &= cov(g, \beta_{my} y) \\
&= cov(g, \beta_{my} \beta_x x) \\
&= cov(g, \beta_{my} \beta_x \beta_g g) \\
&= \beta_{my} \beta_x \beta_g var(g)
\end{aligned}
$$

Part 2

$$
\begin{aligned}
cov(g, \hat{y}_o) &= cov(g, \hat{\beta}_{x_o} x_o) \\
&= cov(g, \hat{\beta}_{x_o} \beta_{mx} x) \\
&= cov(g, \hat{\beta}_{x_o} \beta_{mx} \beta_g g) \\
&= \hat{\beta}_{x_o} \beta_{mx} \beta_g var(g)
\end{aligned}
$$

Simpifying further

$$
\begin{aligned}
\hat{\beta}_{x_o} &= \frac{cov(y_o, x_o)}{var(x_o)} \\
&= \frac{cov(\beta_{my} y, \beta_{mx} x)}{\beta_{mx}^2 var(x) + var(\epsilon_{mx})} \\
&= \frac{\beta_{mx} \beta_{my} cov(y, x)}{\beta_{mx}^2 var(x) + var(\epsilon_{mx})} \\
&= \frac{\beta_{mx} \beta_{my} \beta_x var(x)}{\beta_{mx}^2 var(x) + var(\epsilon_{mx})}
\end{aligned}
$$

which can be substituted back to give

$$
\begin{aligned}
cov(g, \hat{y}_o) &= \frac{\beta_{my} \beta_x \beta_g var(g) \beta_{mx}^2 var(x)}{\beta_{mx}^2 var(x) + var(\epsilon_{mx})} \\
&= \frac{\beta_{mx}^2 var(x)}{\beta_{mx}^2 var(x) + var(\epsilon_{mx})} \times \beta_{my} \beta_x \beta_g var(g)
\end{aligned}
$$

Finally

$$cov(g, y_o - \hat{y}_o) = \beta_{my}\beta_x\beta_g var(g) - \frac{\beta_{mx}^2 var(x)}{\beta_{mx}^2 var(x) + var(\epsilon_m)} \times \beta_{my}\beta_x\beta_g var(g)$$

thus $cov(g, y_o - \hat{y}_o) = 0$ if the measurement imprecision in x_o is $var(\epsilon_m) = 0$. However if there is any imprecision then the condition $cov(g, y_o - \hat{y}_o) = 0$ will not hold.

Appendix 2

The Steiger test is used to infer if the variant g has a direct influence on x or y when it is known that it associates with both, but the direction of causality between x and y is unknown. Assuming the causal direction is $x \to y$, two stage MR is formulated using the following regression models:

$$x = \alpha_1 + \beta_1 g + e_1$$

for the first stage and

$$y = \alpha_2 + \beta_2 \hat{x} + e_2$$

where $\hat{x} = \hat{\alpha}_1 + \hat{\beta}_1 g$. Writing in scale free terms, $\rho_{g,x}$ denotes the correlation between g and the exposure variable x, and it is expected that $\rho_{g,x} > \rho_{g,y}$ because $\rho_{g,y} = \rho_{g,x}\rho_{x,y}$, where $\rho_{x,y}$ is the causal association between x and y (which is likely to be less than 1).

In the presence of measurement error in x and y, however, the Steiger test will instead be assessing the inequality $\rho_{g,x_o} > \rho_{g,y_o}$, which can be simplified:

$$\rho_{g,x_O} > \rho_{g,y_O}$$
$$\rho_{g,x}\rho_{x,x_O} > \rho_{g,y}\rho_{y,y_O}$$
$$\rho_{g,x}\rho_{x,x_O} > \rho_{g,x}\rho_{x,y}\rho_{y,y_O}$$
$$\rho_{x,x_O} > \rho_{x,y}\rho_{y,y_O}$$

In order to assess how reliable the Steiger test is in the presence of measurement imprecision, we can evaluate the range of potential values of measurement error in x and y over which the empirical result from the Steiger test would return the wrong value.

For different values of ρ_{x,x_o}, $\rho_{g,x} = \frac{\rho_{g,x_o}}{\rho_{x,x_o}}$ and $\rho_{g,x_o} \le \rho_{x,x_o} \le 1$. For different values of ρ_{y,y_o}, $\rho_{g,y} = \frac{\rho_{g,y_o}}{\rho_{y,y_o}}$ and $\rho_{g,y_o} \le \rho_{y,x_o} \le 1$.

Where $z = \rho_{g,y} - \rho_{g,x} < 0$ the measurement error range in x_o and y_o will not change the qualitative inference of the direction of causality. To evaluate the reliability, R, of the Steiger test with regards to measurement error, the objective is to compare the weighted proportion of the parameter space that agrees with the inferred direction against the weighted proportion which does not:

$$R = \frac{V_{z \ge 0}}{-V_{z < 0}}$$

If $R = 1$ then the direction of causality is equally probable across the range of possible measurement error values. If $R > 1$ then R times as much of the weighted parameter space favours the inferred direction of causality. V_z, the total volume of the function (Figure 3), can be obtained analytically by solving:

$$V_z = \int_{\rho_{g,x_o}}^{1} \int_{\rho_{g,y_o}}^{1} \frac{\rho_{g,y_o}}{\rho_{y,y_o}} - \frac{\rho_{g,x_o}}{\rho_{x,x_o}} \ d\rho_{y,y_o} d\rho_{x,x_o}$$

$$= \rho_{g,x_o} log(\rho_{g,x_o}) - \rho_{g,y_o} log(\rho_{g,y_o}) + \rho_{g,x_o}\rho_{g,y_o}(log(\rho_{g,y_o}) - log(\rho_{g,x_o}))$$

$V_{z \geq 0}$, the proportion of the volume that lies above the $z = 0$ plane, can also be obtained analytically. The region of this volume is bound by the values of ρ_{x,x_o} and ρ_{y,y_o} where $0 = \rho_{g,y} - \rho_{g,x}$, which can be expanded to $\rho_{y,y_o} = \rho_{g,y_o}\rho_{x,x_o}/\rho_{g,x_o}$. Hence,

$$V_{z \geq 0} = \int_{\rho_{g,x_o}}^{1} \int_{\rho_{g,y_o}}^{\frac{\rho_{g,y_o}\rho_{x,x_o}}{\rho_{g,x_o}}} \frac{\rho_{g,y_o}}{\rho_{y,y_o}} - \frac{\rho_{g,x_o}}{\rho_{x,x_o}} \ d\rho_{y,y_o} d\rho_{x,x_o}$$

$$= 2\rho_{g,x_o}\rho_{g,y_o} - 2\rho_{g,y_o} - \rho_{g,y_o} log(\rho_{g,x_o}) - \rho_{g,x_o}\rho_{g,y_o} log(\rho_{g,x_o})$$

Thus $V_{z<0} = V_z - V_{z \geq 0}$.

References

1. Phillips, A. N. & Davey Smith, G. How independent are 'independent' effects? relative risk estimation when correlated exposures are measured imprecisely. *Journal of Clinical Epidemiology* **44,** 1223–1231 (1991).

2. Davey Smith, G. & Ebrahim, S. Data dredging, bias, or confounding. *BMJ* **325,** (2002).

3. Davey Smith, G. & Ebrahim, S. Mendelian randomization: prospects, potentials, and limitations. *International journal of epidemiology* **33,** 30–42 (2004).

4. Millstein, J., Zhang, B., Zhu, J. & Schadt, E. E. Disentangling molecular relationships with a causal inference test. *BMC genetics* **10,** 23 (2009).

5. Aten, J. E., Fuller, T. F., Lusis, A. J. & Horvath, S. Using genetic markers to orient the edges in quantitative trait networks: the NEO software. *BMC systems biology* **2,** 34 (2008).

6. Waszak, S. M. *et al.* Variation and genetic control of chromatin architecture in humans. *Cell* **162,** 1039–1050 (2015).

7. Houle, D., Pélabon, C., Wagner, G. & Hansen, T. Measurement and meaning in biology. *The Quarterly Review of Biology* **86,** 3–34 (2011).

8. Hernán, M. a & Cole, S. R. Invited Commentary: Causal diagrams and measurement bias. *American journal of epidemiology* **170,** 959–62; discussion 963–4 (2009).

9. Harper, K. N., Peters, B. a & Gamble, M. V. Batch effects and pathway analysis: two potential perils in cancer studies involving DNA methylation array analysis. *Cancer epidemiology, biomarkers & prevention : a publication of the American Association for Cancer Research, cosponsored by the American Society of Preventive Oncology* **22,** 1052–60 (2013).

10. Chen, Y.-a. *et al.* Discovery of cross-reactive probes and polymorphic CpGs in the Illumina Infinium HumanMethylation450 microarray. *Epigenetics : official journal of the DNA Methylation Society* **8,** 203–9 (2013).

11. Houseman, E. A. *et al.* DNA methylation arrays as surrogate measures of cell mixture distribution. *BMC bioinformatics* **13,** 86 (2012).

12. Ahima, R. S. & Lazar, M. A. Physiology. The health risk of obesity–better metrics imperative. *Science (New York, N.Y.)* **341,** 856–8 (2013).

13. Cessie, S. le, Debeij, J., Rosendaal, F. R., Cannegieter, S. C. & Vandenbroucke, J. P. Quantification of bias in direct effects estimates due to different types of measurement error in the mediator. *Epidemiology (Cambridge, Mass.)* **23,** 551–60 (2012).

14. Blakely, T., McKenzie, S. & Carter, K. Misclassification of the mediator matters when estimating indirect effects. *Journal of epidemiology and community health* **67,** 458–66 (2013).

15. Davey Smith, G. & Ebrahim, S. 'Mendelian randomization': can genetic epidemiology contribute to understanding environmental determinants of disease? *International Journal of Epidemiology* **32,** 1–22 (2003).

16. Davey Smith, G. & Hemani, G. Mendelian randomization: genetic anchors for causal inference in epidemiological studies. *Human molecular genetics* **23,** R89—-R98 (2014).

17. Schadt, E. E. *et al.* An integrative genomics approach to infer causal associations between gene expression and disease. *Nature Genetics* **37,** 710–717 (2005).

18. Millstein, J. cit: Causal Inference Test. R package version 1.9. (2016). at <http://cran.r-project.org/package=cit>

19. Koestler, D. C. *et al.* Integrative genomic analysis identifies epigenetic marks that mediate genetic risk

for epithelial ovarian cancer. *BMC medical genomics* **7,** 8 (2014).

20. Liu, Y. *et al.* Epigenome-wide association data implicate DNA methylation as an intermediary of genetic risk in rheumatoid arthritis. *Nature biotechnology* **31,** 142–7 (2013).

21. Yuan, W. *et al.* An integrated epigenomic analysis for type 2 diabetes susceptibility loci in monozygotic twins. *Nature communications* **5,** 5719 (2014).

22. Tang, Y. *et al.* Genotype-based treatment of type 2 diabetes with an α2A-adrenergic receptor antagonist. *Science translational medicine* **6,** 257ra139 (2014).

23. Hong, X. *et al.* Genome-wide association study identifies peanut allergy-specific loci and evidence of epigenetic mediation in US children. *Nature communications* **6,** 6304 (2015).

24. Bowden, J., Davey Smith, G. & Burgess, S. Mendelian randomization with invalid instruments: effect estimation and bias detection through Egger regression. *International Journal of Epidemiology* **In press,** (2015).

25. Bowden, J., Davey Smith, G., Haycock, P. C. & Burgess, S. Consistent Estimation in Mendelian Randomization with Some Invalid Instruments Using a Weighted Median Estimator. *Genetic Epidemiology* **40,** 304–314 (2016).

26. Hemani, G. *et al.* MR-Base: a platform for systematic causal inference across the phenome using billions of genetic associations. *BioRxiv* **10.1101/07,** (2016).

27. Claussnitzer, M. *et al.* FTO Obesity Variant Circuitry and Adipocyte Browning in Humans. *The New England journal of medicine* **373,** 895–907 (2015).

28. Shakhbazov, K. *et al.* Shared genetic control of expression and methylation in peripheral blood. *BMC genomics* **17,** 278 (2016).

29. Pierce, B. L. & Burgess, S. Efficient design for Mendelian randomization studies: subsample and 2-sample instrumental variable estimators. *American journal of epidemiology* **178,** 1177–84 (2013).

30. Ashenfelter, O. & Krueger, A. B. Estimates of the Economic Return to Schooling from a New Sample of Twins. *The American Economic Review* **84,** 1157–1173 (1994).

31. Nagarajan, R. & Scutari, M. Impact of noise on molecular network inference. *PloS one* **8,** e80735 (2013).

32. Shpitser, I., VanderWeele, T. & Robins, J. On the validity of covariate adjustment for estimating causal effects. *Proceedings of the Twenty Sixth Conference on Uncertainty in Artificial Intelligence (UAI-10)* 527–536 (2010).

33. Wang, L. & Michoel, T. Detection of regulator genes and eQTLs in gene networks. *arXiv* **arXiv:1512,** (2015).

34. Lagani, V., Triantafillou, S., Ball, G., Tegner, J. & Tsamardinos, I. in *Uncertainty in biology: A computational modeling approach* 47 (Springer, 2015). at <https://books.google.com/books?id=8SLUCgAAQBAJ{\&}pgis=1>

35. R Core Team. R: A Language and Environment for Statistical Computing. (2015). at <https://www.r-project.org/>

36. Sterne, J. A. C., Cox, D. R. & Smith, G. D. Sifting the evidence—what's wrong with significance tests?Another comment on the role of statistical methods. *BMJ* **322,** (2001).

37. Henningsen, A. & Hamann, J. D. systemfit : A Package for Estimating Systems of Simultaneous Equations in R. *Journal of Statistical Software* **23,** 1–40 (2007).

38. Steiger, J. H. Tests for comparing elements of a correlation matrix. *Psychological Bulletin* **87,** 245–251 (1980).

39. Revelle, W. psych: Procedures for Psychological, Psychometric, and Personality Research. (2015). at

<http://cran.r-project.org/package=psych>

40. Bird, A. DNA methylation patterns and epigenetic memory. *Genes & development* **16,** 6–21 (2002).

41. Fewell, Z., Davey Smith, G. & Sterne, J. A. C. The impact of residual and unmeasured confounding in epidemiologic studies: a simulation study. *American journal of epidemiology* **166,** 646–55 (2007).

42. Bose, M. *et al.* Evaluation of microarray-based DNA methylation measurement using technical replicates: the Atherosclerosis Risk In Communities (ARIC) Study. *BMC Bioinformatics* **15,** 312 (2014).

43. Bryant, P. A. *et al.* Technical Variability Is Greater than Biological Variability in a Microarray Experiment but Both Are Outweighed by Changes Induced by Stimulation. *PLoS ONE* **6,** e19556 (2011).

44. Richmond, R. C. *et al.* Assessing Causality in the Association between Child Adiposity and Physical Activity Levels: A Mendelian Randomization Analysis. *PLoS Medicine* **11,** e1001618 (2014).

45. Relton, C. L. & Davey Smith, G. Two-step epigenetic Mendelian randomization: a strategy for establishing the causal role of epigenetic processes in pathways to disease. *International journal of epidemiology* **41,** 161–76 (2012).

46. Burgess, S., Freitag, D. F., Khan, H., Gorman, D. N. & Thompson, S. G. Using multivariable Mendelian randomization to disentangle the causal effects of lipid fractions. *PloS one* **9,** e108891 (2014).

47. Varbo, A. *et al.* Remnant cholesterol, low-density lipoprotein cholesterol, and blood pressure as mediators from obesity to ischemic heart disease. *Circulation research* **116,** 665–73 (2015).

48. Burgess, S., Daniel, R. M., Butterworth, A. S. & Thompson, S. G. Network Mendelian randomization: using genetic variants as instrumental variables to investigate mediation in causal pathways. *International journal of epidemiology* **44,** 484–95 (2015).